

Bios 6301: Assignment 7

Jonathan Lifferth

Due Thursday, 02 November, 1:00 PM

$5^{n=\text{day}}$ points taken off for each day late.

40 points total.

Submit a single knitr file (named `homework7.rmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework7.rmd` or include author name may result in 5 points taken off.

Question 1

21 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {  
  if(exists(".Random.seed", envir = .GlobalEnv)) {  
    save.seed <- get(".Random.seed", envir = .GlobalEnv)  
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))  
  } else {  
    on.exit(rm(".Random.seed", envir = .GlobalEnv))  
  }  
  set.seed(n)  
  subj <- ceiling(n / 10)  
  id <- sample(subj, n, replace=TRUE)  
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))  
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')  
  mu <- runif(subj, 4, 10)  
  a1c <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)  
  data.frame(id, dt, a1c)  
}  
x <- genData(500)
```

Perform the following manipulations: (3 points each)

1. Order the data set by `id` and `dt`.

```
head(x)
```

```
##      id              dt      a1c
## 1 39 2001-11-30 00:36:10 8.758993
## 2 27 2000-07-20 17:05:06 8.233413
## 3 47 2001-02-06 18:40:44 3.950582
## 4 50 2001-04-05 01:55:44 6.707085
## 5 41 2003-08-05 18:30:14 4.047338
## 6 31 2000-12-16 10:38:26 9.357902
```

```
x<-x[order(x[, 'id'], x[, 'dt']),]
head(x)
```

```
##      id              dt      a1c
## 32  1 2001-05-08 16:22:52 7.309995
## 268 1 2001-06-17 22:42:23 8.310721
## 201 1 2001-08-17 16:51:46 6.548845
## 285 1 2001-12-14 14:50:29 5.985275
## 194 1 2002-08-19 13:51:47 6.011547
## 304 1 2003-03-22 03:51:36 7.243858
```

2. For each `id`, determine if there is more than a one year gap in between observations. Add a new row at the one year mark, with the `a1c` value set to missing. A two year gap would require two new rows, and so forth.

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:base':
##
##      date, intersect, setdiff, union
```

```
x <- genData(500)
x_ordered<-x[order(x[, 'id'], x[, 'dt']),]
rownames(x_ordered) <- NULL

x <- x[order(x[, 'id'], x[, 'dt']),]
rownames(x) <- NULL
gap_years <- data.frame()

nrow(x)
```

```
## [1] 500
```

```
for (i in 1:(nrow(x)-1)) {
  current_id <- x[i, 'id']
  next_id <- x[i+1, 'id']
```

```

# only proceed if we're evaluating the same id
if (next_id == current_id) {
  # what is the time gap between the current row and the next row?
  delta <- abs(difftime(x[i+1,'dt'], x[i,'dt'], units = 'days'))[[1]]
  num_years <- trunc(delta / 365)

  # add new rows for each observation year gap
  if (num_years > 0) {
    for (j in 1:num_years) {
      new_row <- data.frame(id = current_id, dt = x[i,'dt']+years(j), a1c=NA)
      # x <- rbind(x[1:i,], new_row, x[i+1:nrow(x),])
      gap_years <- rbind(gap_years, new_row)
    }
  }
}

```

```
nrow(gap_years)
```

```
## [1] 45
```

```

x2 <- rbind(x, gap_years)
x2<-x2[order(x2[, 'id'], x2[, 'dt']),]
rownames(x2) <- NULL

head(x2)

```

```

##   id          dt      a1c
## 1  1 2001-05-08 16:22:52 7.309995
## 2  1 2001-06-17 22:42:23 8.310721
## 3  1 2001-08-17 16:51:46 6.548845
## 4  1 2001-12-14 14:50:29 5.985275
## 5  1 2002-08-19 13:51:47 6.011547
## 6  1 2003-03-22 03:51:36 7.243858

```

3. Create a new column `visit`. For each `id`, add the visit number. This should be 1 to `n` where `n` is the number of observations for an individual. This should include the observations created with missing `a1c` values.

```

x2$visit <- 0
count <- 1 # this is the default value for visit number
for (i in 1:nrow(x2)) {

  # we need to exclude the very first row, because it can't be compared to any previous row ids
  if (i > 1) {
    if (x2[i, 'id'] == x2[i-1, 'id']) {
      count <- count + 1
      x2[i, 'visit'] <- count
    }
    else { # this indicates the first row for a new id

```

```

    count <- 1
    x2[i, 'visit'] <- count
  }

} else { # this condition should only apply to the very first row in the data.frame
  x2[i, 'visit'] <- 1
}
}

```

4. For each id, replace missing values with the mean a1c value for that individual.

```

for (i in 1:nrow(x2)) {
  if (is.na(x2[i, 'a1c'])) {
    x2[i, 'a1c'] <- mean(subset(x2, id == x2[i, 'id'])$a1c, na.rm=TRUE)
  }
}

```

5. Print mean a1c for each id.

```

for (i in unique(x2$id)) {
  print(paste('id: ', i, ' ', mean(subset(x2, id == i)$a1c, na.rm=TRUE)))
}

```

```

## [1] "id: 1      6.65444426795186"
## [1] "id: 2      9.78913246074151"
## [1] "id: 3      6.95182045895334"
## [1] "id: 4      8.19198450682839"
## [1] "id: 5      9.42969414135007"
## [1] "id: 6      7.13344348656912"
## [1] "id: 7      7.87913801432509"
## [1] "id: 8      6.24406099245875"
## [1] "id: 9      4.42052304020483"
## [1] "id: 10     6.02836978936866"
## [1] "id: 11     4.83827911476455"
## [1] "id: 12     6.69118108424096"
## [1] "id: 13     8.50463215686808"
## [1] "id: 14     9.12296781957672"
## [1] "id: 15     6.73709205512209"
## [1] "id: 16     7.42024462564604"
## [1] "id: 17     6.54632858730216"
## [1] "id: 18     6.1513112940644"
## [1] "id: 19     8.62803745758515"
## [1] "id: 20     8.92351824057672"
## [1] "id: 21     5.444430006372"
## [1] "id: 22     5.76393126014759"
## [1] "id: 23     6.35111217834161"
## [1] "id: 24     9.37752492553745"
## [1] "id: 25     5.05809652490457"
## [1] "id: 26     8.69207762927627"
## [1] "id: 27     7.37183147872539"
## [1] "id: 28     4.24346852483802"
## [1] "id: 29     6.34525429737664"

```

```
## [1] "id: 30      4.13579498572139"
## [1] "id: 31      8.67062198152496"
## [1] "id: 32      5.1301670704902"
## [1] "id: 33      6.52815306924961"
## [1] "id: 34      8.44503021368734"
## [1] "id: 35      3.83219482233089"
## [1] "id: 36      9.51460255980355"
## [1] "id: 37      8.61260794411042"
## [1] "id: 38      10.160772908825"
## [1] "id: 39      8.97669727861485"
## [1] "id: 40      7.58323173368407"
## [1] "id: 41      3.8043252144796"
## [1] "id: 42      6.78716991115953"
## [1] "id: 43      5.65423470328969"
## [1] "id: 44      5.61328261848045"
## [1] "id: 45      8.8766234785112"
## [1] "id: 46      7.4858240579994"
## [1] "id: 47      4.75213278333204"
## [1] "id: 48      7.41545866940117"
## [1] "id: 49      5.56280902415056"
## [1] "id: 50      4.97028797276639"
```

6. Print total number of visits for each id.

```
for (i in unique(x2$id)) {
  print(paste('id: ', i, ' ', max(subset(x2,id == i)$visit)))
}
```

```
## [1] "id: 1      7"
## [1] "id: 2     16"
## [1] "id: 3     13"
## [1] "id: 4      9"
## [1] "id: 5     14"
## [1] "id: 6     11"
## [1] "id: 7      7"
## [1] "id: 8     12"
## [1] "id: 9     15"
## [1] "id: 10     8"
## [1] "id: 11     12"
## [1] "id: 12     12"
## [1] "id: 13      9"
## [1] "id: 14     12"
## [1] "id: 15     10"
## [1] "id: 16      8"
## [1] "id: 17     10"
## [1] "id: 18     14"
## [1] "id: 19     10"
## [1] "id: 20     11"
## [1] "id: 21     13"
## [1] "id: 22     12"
## [1] "id: 23     10"
## [1] "id: 24     12"
## [1] "id: 25     16"
```

```
## [1] "id: 26      11"
## [1] "id: 27      10"
## [1] "id: 28      15"
## [1] "id: 29       3"
## [1] "id: 30      13"
## [1] "id: 31      11"
## [1] "id: 32       9"
## [1] "id: 33      12"
## [1] "id: 34      12"
## [1] "id: 35      11"
## [1] "id: 36      10"
## [1] "id: 37       8"
## [1] "id: 38      14"
## [1] "id: 39      14"
## [1] "id: 40      11"
## [1] "id: 41      14"
## [1] "id: 42      11"
## [1] "id: 43       8"
## [1] "id: 44      12"
## [1] "id: 45       6"
## [1] "id: 46      12"
## [1] "id: 47      10"
## [1] "id: 48       5"
## [1] "id: 49      11"
## [1] "id: 50       9"
```

7. Print the observations for id = 15.

```
subset(x2, id==15)
```

```
##      id      dt      a1c visit
## 158 15 2000-10-21 01:08:17 7.401322    1
## 159 15 2001-08-08 14:23:08 5.896318    2
## 160 15 2001-08-15 07:03:29 7.457722    3
## 161 15 2002-03-15 21:23:10 5.330917    4
## 162 15 2002-04-14 09:08:25 6.484003    5
## 163 15 2002-10-10 18:27:43 8.139101    6
## 164 15 2003-02-19 12:58:53 6.446557    7
## 165 15 2003-03-02 06:58:10 7.432291    8
## 166 15 2003-06-30 07:20:49 7.113792    9
## 167 15 2004-01-22 20:30:42 5.668897   10
```

Question 2

16 points

Install the `lexicon` package. Load the `sw_fry_1000` vector, which contains 1,000 common words.

```
#install.packages('lexicon')
library('lexicon')
```

```
data('sw_fry_1000', package = 'lexicon')
head(sw_fry_1000)
```

```
## [1] "the" "of" "to" "and" "a" "in"
```

```
length(sw_fry_1000)
```

```
## [1] 1000
```

1. Remove all non-alphabetical characters and make all characters lowercase. Save the result as a.

```
sw_fry_1000[which(grepl("[^[:alpha:]]", sw_fry_1000, ignore.case = TRUE))]
```

```
## [1] "don't" "won't"
```

```
a <- gsub("[^[:alpha:]]", "", sw_fry_1000)
head(a)
```

```
## [1] "the" "of" "to" "and" "a" "in"
```

```
a <- sapply(a, tolower)
```

Use vector a for the following questions. (2 points each)

2. How many words contain the string “ar”?

```
sum(grepl("ar", a))
```

```
## [1] 64
```

3. Find a six-letter word that starts with “l” and ends with “r”.

```
# Create a regular expression to match the letter "l" at the beginning of the word
l_regex <- "^l"
```

```
# Create a regular expression to match the letter "r" at the end of the word
r_regex <- "r$"
```

```
indices <- which(grepl(l_regex, a) & grepl(r_regex, a))
```

```
print(length(a[indices[1]]))
```

```
## [1] 1
```

```
print(a[indices[1]])
```

```
## letter
## "letter"
```

4. Return all words that start with “col” or end with “eck”.

```
a[which(grepl("^col", a) | grepl("eck$", a))]
```

```
##      color      cold      check      collect      colony      column      neck
## "color"      "cold"      "check" "collect"      "colony"      "column"      "neck"
```

5. Find the number of words that contain 4 or more adjacent consonants. Assume “y” is always a consonant.

```
print(a[which(grepl('[bcdfghjklmnpqrstvwxyz]{4,}', a))])
```

```
##      country      system      syllable      length      instrument      industry
## "country"      "system"      "syllable"      "length" "instrument"      "industry"
##      symbol      supply
## "symbol"      "supply"
```

```
print(length(which(grepl('[bcdfghjklmnpqrstvwxyz]{4,}', a))))
```

```
## [1] 8
```

6. Return all words with a “q” that isn’t followed by a “ui”.

```
grep("\\b\\w*q(?!ui)\\w*\\b", a, value = TRUE, perl = TRUE)
```

```
##      question      equate      square      equal      quart      quotient
## "question"      "equate"      "square"      "equal"      "quart"      "quotient"
```

7. Find all words that contain a “k” followed by another letter. Run the `table` command on the first character following the first “k” of each word.

```
k_words <- grep("\\b\\w*k\\w+\\b", a, value=TRUE)
k_words
```

```
##      like      make      know      take      kind      keep      knew      king
## "like"      "make"      "know"      "take"      "kind"      "keep"      "knew"      "king"
##      sky      kept      broke      kill      lake      key      skin      spoke
## "sky"      "kept"      "broke"      "kill"      "lake"      "key"      "skin"      "spoke"
##      skill      market
## "skill"      "market"
```

```
first_chars <- character(0)
```

```
for (word in k_words) {
  match <- regexpr("k.", word)
  if (match != -1) {
    first_char <- substr(word, match + 1, match + 1)
    first_chars <- c(first_chars, first_char)
  }
}
```



```
# Use the table function to count the occurrences of each unique first character
char_count <- table(first_chars)
```

```
# Print the character count
print(char_count)
```

```
## first_chars
## e i n y
## 10 5 2 1
```

8. Remove all vowels. How many character strings are found exactly once?

```
remove_vowels <- function(string) {
  return(gsub("[aeiouAEIOU]", "", string, ignore.case = TRUE))
}
no_vowels <- lapply(a, remove_vowels)
length(unique(no_vowels))
```

```
## [1] 741
```

Question 3

3 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula.

```
url <- "https://github.com/couthcommander/Bios6301/raw/master/datasets/haart.csv"
url <- "https://github.com/couthcommander/Bios6301/blob/main/datasets/haart.csv"
```

```
haart_df <- read.csv(url)[,c('death', 'weight', 'hemoglobin', 'cd4baseline')]
```

```
## Warning in read.table(file = file, header = header, sep = sep,
## quote = quote, : incomplete final line found by readTableHeader on
## 'https://github.com/couthcommander/Bios6301/blob/main/datasets/haart.csv'
```

```
# the urls weren't loading so I used a copy of the dataset that I already had downloaded
```

```
haart_df <- read.csv('/Users/jonathanlifferth/R_projects/Bios6301/datasets/haart.csv')[,c('death', 'weight', 'hemoglobin', 'cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(predvars, data, env): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

```
# debug(myfun(dat=haart_df, response=death))
```

Debugging revealed the following error: "Error in eval(predvars, data, env) : object 'death' not found" Meaning that the `death` column in `haart_df` is not accessible in the scope of this function.

5 bonus points

Create a working function.

`eval()` can't find `death` in the current scope/environment

```
myfun2 <- function(dat, response) {
  form <- paste0(substitute(response), " ~ .")
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
myfun2(dat=haart_df, response=death)
```

```
##               Estimate Std. Error  z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535  2.915069 0.0035561039
## weight      -0.046210552 0.022556001 -2.048703 0.0404911395
## hemoglobin   -0.350642786 0.105064078 -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959  1.154872 0.2481427160
```